



## Part 3: The Entity-Relationship Model

**Database System Concepts, 7<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Design Phases

- Initial phase -- characterize fully the data needs of the prospective database users
- Second phase -- choosing a data model
  - Applying the concepts of the chosen data model
  - Translating these requirements into a conceptual schema of the database
  - A fully developed conceptual schema indicates the functional requirements of the enterprise
    - Describe the kinds of operations that will be performed on the data



# Design Phases

- Final Phase -- Moving from an abstract data model to the implementation of the database
  - Logical Design – Deciding on the database schema. Database design requires that we find a “good” collection of relation schemas
    - Business decision – What attributes should we record in the database?
    - Computing decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
  - Physical Design – Deciding on the physical layout of the database



# Design Alternatives

- In designing a database schema, we must ensure that we avoid two major pitfalls:
  - Redundancy: a bad design may result in repeated information.
    - Redundant representation of information may lead to data inconsistency among the various copies of information
  - Incompleteness: a bad design may make certain aspects of the enterprise difficult or impossible to model
- Avoiding bad designs is not enough. There may be a large number of good designs from which we must choose



# Design Approaches

- Entity-Relationship Model
  - Models an enterprise as a collection of *entities* and *relationships*
    - Entity: a “thing” or “object” in the enterprise that is distinguishable from other objects
      - Described by a set of *attributes*
    - Relationship: an association among several entities
  - Represented diagrammatically by an *entity-relationship diagram*
- Normalization Theory
  - Formalize what designs are bad, and test for them



# Outline of the ER Model



# ER model -- Database Modeling

- The ER data model was developed to facilitate database design by allowing specification of an **enterprise schema** that represents the overall logical structure of a database
- The ER data model employs three basic concepts:
  - entity sets
  - relationship sets
  - attributes
- The ER model also has an associated diagrammatic representation, the **ER diagram**, which can express the overall logical structure of a database graphically



# Entity Sets



- An **entity** is an object that exists and is distinguishable from other objects.
  - Example: specific person, company, event, plant
- An **entity set** is a set of entities of the same type that share the same properties
  - Example: set of all persons, companies, products
- An entity is represented by a set of attributes; i.e., descriptive properties possessed by all members of an entity set
  - Example:  
 $instructor = (ID, name, salary)$   
 $course = (course\_id, title, credits)$
- A subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set





# Representing Entity sets in ER Diagram

- Entity sets can be represented graphically as follows:
  - Rectangles represent entity sets.
  - Attributes listed inside entity rectangle
  - Underline indicates primary key attributes





# Relationship Sets

- A **relationship** is an association among several entities

Example:

44553 (Peltier)	<u>advisor</u>	22222 ( <u>Einstein</u> )
student entity	relationship set	instructor entity

- A **relationship set** is a mathematical relation among entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) : e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where  $(e_1, e_2, \dots, e_n)$  is a relationship

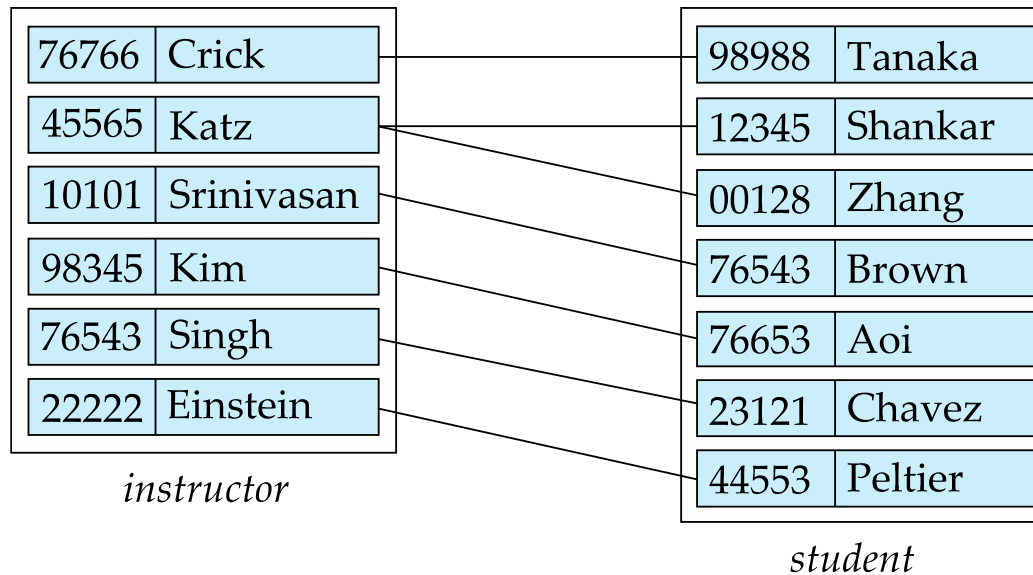
- Example:

$$(44553, 22222) \in \text{advisor}$$



# Relationship Sets


- Example: we define the relationship set *advisor* to denote the associations between students and the instructors who act as their advisors
- Pictorially, we draw a line between related entities





# Representing Relationship Sets via ER Diagrams

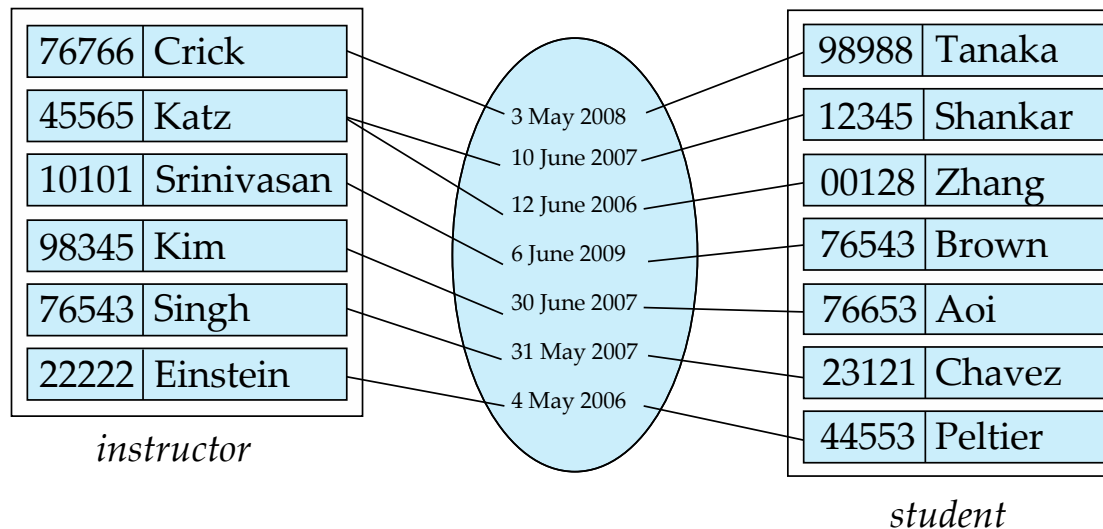
- Diamonds represent relationship sets






# Relationship Sets

- An attribute can also be associated with a relationship set, sometimes called a **descriptive attribute**
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date* which tracks when the student started being associated with the advisor





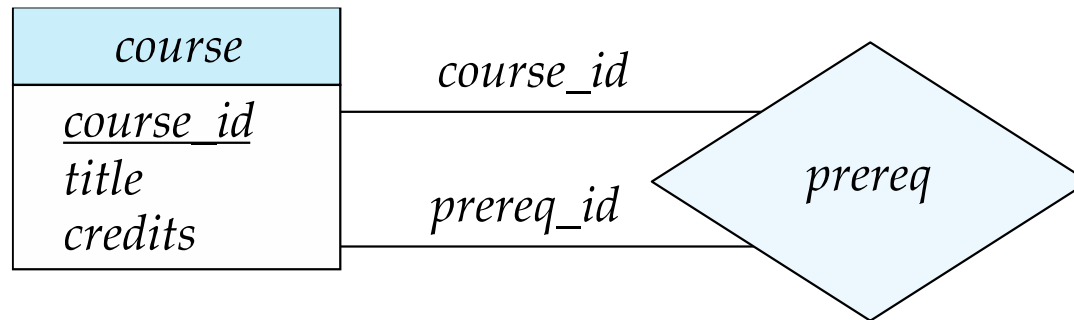
# Relationship Sets with Attributes





# Roles


- Entity sets of a relationship need not be distinct
  - Each occurrence of an entity set plays a “role” in the relationship
- The labels “*course\_id*” and “*prereq\_id*” are called **roles**





# Degree of a Relationship Set

- Binary relationship
  - involve two entity sets (or degree two)
  - most relationship sets in a database system are binary
- Relationships between more than two entity sets are less common
  - Example: *students* work on research *projects* under the guidance of an *instructor*
  - relationship *proj\_guide* is a ternary relationship between *instructor*, *student*, and *project*







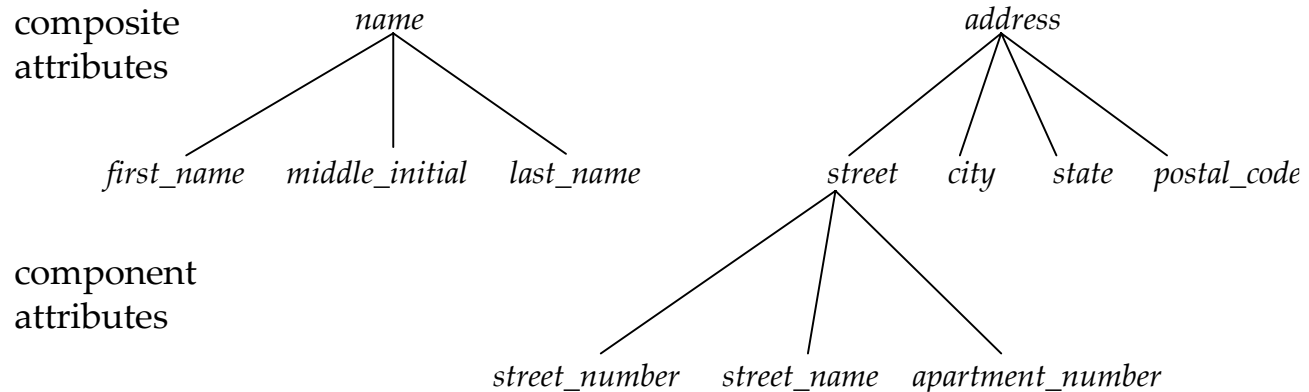
# Complex Attributes

- Attribute types:
  - **Simple** and **composite** attributes
    - Name may consist of first name and last name
  - **Single-valued** and **multivalued** attributes
    - Example: multivalued attribute: phone\_numbers
  - **Derived** attributes
    - Can be computed from other attributes
    - Example: age, given date\_of\_birth
- **Domain** – the set of permitted values for each attribute



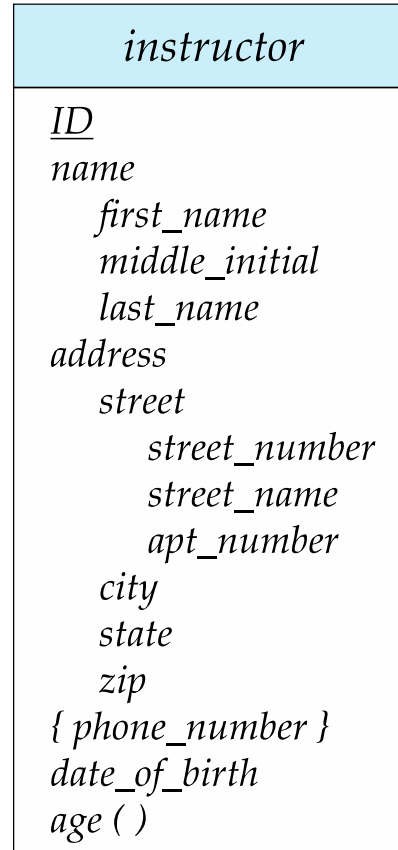
# Composite Attributes

- Composite attributes allow us to divide attributes into subparts (other attributes).





# Representing Complex Attributes in ER Diagram






# Mapping Cardinality Constraints


- Express the number of entities to which another entity can be associated via a relationship set
- Most useful in describing binary relationship sets
- For a binary relationship set the mapping cardinality must be one of the following types:
  - One-to-one
  - One-to-many
  - Many-to-one
  - Many-to-many



# Mapping Cardinalities



One-to-one




One-to-many


Note: Some elements in  $A$  and  $B$  may not be mapped to any elements in the other set



# Mapping Cardinalities



Many-to-one




Many-to-many

Note: Some elements in  $A$  and  $B$  may not be mapped to any elements in the other set



# Representing Cardinality Constraints in ER Diagram


- We express cardinality constraints by drawing either a directed line ( $\rightarrow$ ), signifying “one,” or an undirected line ( $—$ ), signifying “many,” between the relationship set and the entity set
- One-to-one relationship between an *instructor* and a *student*





# One-to-Many Relationship

- one-to-many relationship between an *instructor* and a *student*
  - an instructor is associated with several (including 0) students via *advisor*
  - a student is associated with at most one instructor via *advisor*








# Many-to-One Relationship


- In a many-to-one relationship between an *instructor* and a *student*,
  - an *instructor* is associated with at most one *student* via *advisor*,
  - and a *student* is associated with several (including 0) *instructors* via *advisor*





# Many-to-Many Relationship


- An instructor is associated with several (possibly 0) students via *advisor*
- A student is associated with several (possibly 0) instructors via *advisor*





# Total and Partial Participation

- **Total participation** (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set




participation of *student* in *advisor* relation is total

- every *student* must have an associated instructor
- **Partial participation:** some entities may not participate in any relationship in the relationship set
  - Example: participation of *instructor* in *advisor* is partial



# Notation for Expressing More Complex Constraints

- A line may have an associated minimum and maximum cardinality, shown in the form  $l..h$ , where  $l$  is the minimum and  $h$  the maximum cardinality
  - A minimum value of 1 indicates total participation
  - A maximum value of 1 indicates that the entity participates in at most one relationship
  - A maximum value of \* indicates no limit
- Example



- Instructor can advise 0 or more students. A student must have 1 advisor; cannot have multiple advisors



# Primary Key

- Primary keys provide a way to specify how entities and relations are distinguished. We will consider:
  - Entity sets
  - Relationship sets
  - Weak entity sets



# Primary key for Entity Sets

- By definition, individual entities are distinct
- From a database perspective, the differences among them must be expressed in terms of their attributes
- The values of the attribute values of an entity must be such that they can uniquely identify the entity
  - No two entities in an entity set are allowed to have exactly the same value for all attributes
- A key for an entity is a set of attributes that suffice to distinguish entities from each other



# Primary Key for Relationship Sets

- To distinguish among the various relationships of a relationship set we use the individual primary keys of the entities in the relationship set
  - Let  $R$  be a relationship set involving entity sets  $E_1, E_2, \dots, E_n$
  - The primary key for  $R$  consists of the union of the primary keys of entity sets  $E_1, E_2, \dots, E_n$
- Example: relationship set *advisor*
  - The primary key consists of *instructor.ID* and *student.ID*



# Choice of Primary key for Binary Relationship

- Many-to-Many relationships. The union of the primary keys is a minimal superkey and is chosen as the primary key
- One-to-Many relationships . The primary key of the “Many” side is a minimal superkey and is used as the primary key
  - If a department has many instructors, then the primary key of the instructor-department relationship is simply the primary key of instructor
- Many-to-one relationships. The primary key of the “Many” side is a minimal superkey and is used as the primary key
  - If the relationship is many-to-one from *student* to *instructor*—that is, each student can have at most one advisor—then the primary key of *advisor* is simply the primary key of *student*
- One-to-one relationships. The primary key of either one of the participating entity sets forms a minimal superkey, and either one can be chosen as the primary key





# Weak Entity Sets

- Consider a *section* entity, which is uniquely identified by a *course\_id*, *semester*, *year*, and *sec\_id*
- Clearly, section entities are related to course entities. Suppose we create a relationship set *sec\_course* between entity sets *section* and *course*
- Note that some information in *sec\_course* is redundant, since *section* already has an attribute *course\_id*, which identifies the course with which the section is related
- One option to deal with this redundancy is to get rid of the relationship *sec\_course*; however, by doing so the relationship between *section* and *course* becomes implicit in an attribute, which is not desirable



# Weak Entity Sets

- An alternative way to deal with this redundancy is to not store the attribute *course\_id* in the *section* entity and to only store the remaining attributes *section\_id*, *year*, and *semester*
  - However, the entity set *section* then does not have enough attributes to identify a particular *section* entity uniquely
- To deal with this problem, we treat the relationship *sec\_course* as a special relationship that provides extra information, in this case the *course\_id*, required to identify *section* entities uniquely
- A **weak entity set** is one whose existence is dependent on another entity, called its **identifying entity**
- Instead of associating a primary key with a weak entity, we use the identifying entity, along with extra attributes, called **discriminator attributes**, to uniquely identify a weak entity



# Weak Entity Sets

- An entity set that is not a weak entity set is termed a **strong entity set**.
- Every weak entity must be associated with an identifying entity; that is, the weak entity set is said to be **existence dependent** on the identifying entity set
- The identifying entity set is said to **own** the weak entity set that it identifies
- The relationship associating the weak entity set with the identifying entity set is called the **identifying relationship**
- Note that the relational schema we eventually create from the entity set *section* does have the attribute *course\_id*, even though we have dropped the attribute *course\_id* from the entity set *section*



# Expressing Weak Entity Sets


- In E-R diagrams, a weak entity set is depicted via a double rectangle
- We underline the discriminator of a weak entity set with a dashed line
- The relationship set connecting the weak entity set to the identifying strong entity set is depicted by a double diamond
- In general, a weak entity must have total participation in its identifying relationship set, and the relationship is many-to-one towards the identifying entity set
- Primary key for *section* – (*course\_id*, *sec\_id*, *semester*, *year*)





# Redundant Attributes

- Suppose we have entity sets:
  - *instructor*, with attributes: *ID*, *name*, *dept\_name*, *salary*
  - *department*, with attributes: *dept\_name*, *building*, *budget*
- We model the fact that each instructor has an associated department using a relationship set *inst\_dept*
- The attribute *dept\_name* in *instructor* replicates information present in the relationship and is therefore redundant
- When converting back to tables, in some cases (e.g., when each instructor has exactly one associated department) the attribute gets reintroduced, as we will see later





# Reduction to Relation Schemas



# Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas* that represent the contents of the database
- A database which conforms to an E-R diagram can be represented by a collection of schemas
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set
- Each schema has a number of columns (generally corresponding to attributes), which have unique names



# Representing Entity Sets


- A strong entity set reduces to a schema with the same attributes

*student(ID, name, tot\_cred)*

- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set

*section (course\_id, sec\_id, semester, year)*

- Example







# Representation of Entity Sets with Composite Attributes

<i>instructor</i>
<u>ID</u>
name
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
address
street
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
city
state
zip
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ( )

- Composite attributes are flattened out by creating a separate attribute for each component attribute
  - Example: given entity set *instructor* with composite attribute *name* with component attributes *first\_name* and *last\_name* the schema corresponding to the entity set has two attributes *name\_first\_name* and *name\_last\_name*
    - Prefix omitted if there is no ambiguity (*name\_first\_name* could be *first\_name*)
- Ignoring multivalued attributes, extended instructor schema is
  - *instructor*(ID, *first\_name*, *middle\_initial*, *last\_name*, *street\_number*, *street\_name*, *apt\_number*, *city*, *state*, *zip\_code*, *date\_of\_birth*)



# Representation of Entity Sets with Multivalued Attributes


- A multivalued attribute  $M$  of an entity  $E$  is represented by a separate schema  $EM$
- Schema  $EM$  has attributes corresponding to the primary key of  $E$  and an attribute corresponding to multivalued attribute  $M$
- Example: Multivalued attribute *phone\_number* of *instructor* is represented by a schema:  
$$inst\_phone = ( \underline{ID}, \underline{phone\_number} )$$
- Each value of the multivalued attribute maps to a separate tuple of the relation on schema  $EM$ 
  - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:  
(22222, 456-7890) and (22222, 123-4567)



# Representing Relationship Sets

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*

*advisor* = (*s id*, *i id*)



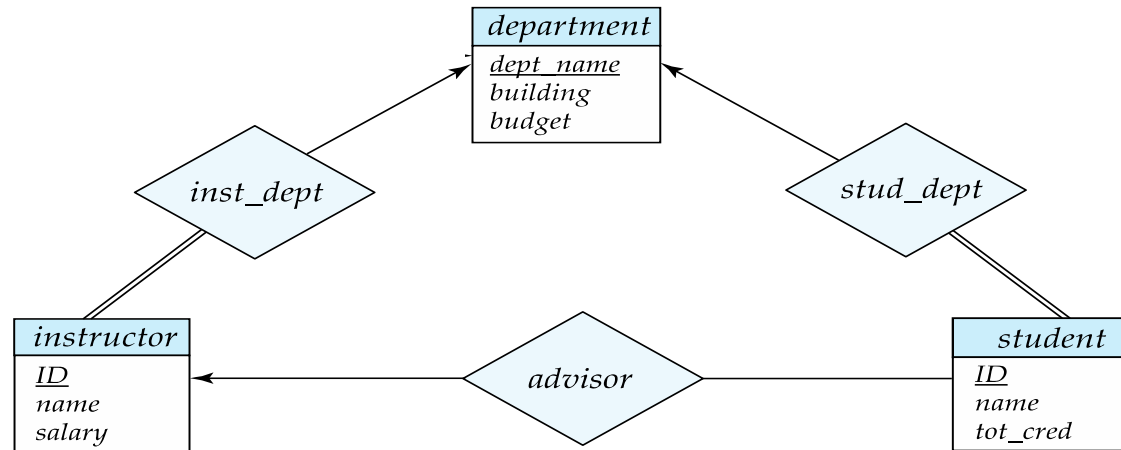


# Combination of Schemas

- Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side
  - Example: Instead of creating a schema for relationship set *inst\_dept*, add an attribute *dept\_name* to the schema arising from entity set *instructor*
- For one-to-one relationship sets, either side can be chosen to act as the “many” side
  - That is, an extra attribute can be added to either of the tables corresponding to the two entity sets



# Combination of Schemas



## *Inst\_dept*

- The schema *inst\_dept* can be combined with the *instructor* schema - the resulting *instructor* schema consists of the attributes {*ID*, *name*, *dept\_name*, *salary*}

## *stud\_dept*

- The schema *stud\_dept* can be combined with the *student* schema - the resulting *student* schema consists of the attributes {*ID*, *name*, *dept\_name*, *tot\_cred*}



# Combination of Schemas

- As we saw, a weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set - the resulting *section* schema consists of the attributes  
 $\{course\_id, sec\_id, semester, year\}$
- *sec\_class*. The schema *sec\_class* can be combined with the *section* schema obtained above - the resulting *section* schema now consists of the attributes  
 $\{course\_id, sec\_id, semester, year, building, room\_number\}$
- The classroom table is still needed to store classroom info (such as capacity), which is necessary, since even if there is no class held in a classroom, its capacity info is still available

