



# Part 11: B-Trees

**Database System Concepts, 7<sup>th</sup> Ed.**


©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use




# Search Trees

- Search tree used to guide search for a record





# Search Trees





# B-Trees

- Provide multi-level access structure
- Tree is always balanced
- Space wasted by deletion never becomes excessive
  - **Each node is at least half-full**
- A B-tree can store values in non-leaf and leaf nodes
- Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children; i.e.,  $n$  signifies the maximum number of children, and  $n$  is called the **order** of the tree



# B<sup>+</sup>-Tree Index Files

A B<sup>+</sup>-tree is a rooted tree satisfying the following properties:

- All paths from root to leaf are of the same length
- Non-leaf nodes are called **internal nodes**
- Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children
- Note that the number of values is one less than the number of children; i.e., between  $\lceil n/2 \rceil - 1$  and  $n - 1$  values, or equivalently, between  $\lfloor (n-1)/2 \rfloor$  and  $n - 1$  values
- The above equivalence can be seen by noting that  **$\lceil n/2 \rceil - 1 = \lfloor (n-1)/2 \rfloor$** ; this is so because
  - for  $n$  even,  $\lceil n/2 \rceil - 1 = (n/2) - 1$ , and  $\lfloor (n-1)/2 \rfloor = \lfloor (n/2) - 1/2 \rfloor = (n/2) - 1$ ;  
thus  $\lceil n/2 \rceil - 1 = \lfloor (n-1)/2 \rfloor = (n/2) - 1$
  - for  $n$  odd,  $\lceil n/2 \rceil - 1 = (n+1)/2 - 1 = (n-1)/2$ , and  $\lfloor (n-1)/2 \rfloor = (n-1)/2$ ;  
thus  $\lceil n/2 \rceil - 1 = \lfloor (n-1)/2 \rfloor = (n-1)/2$

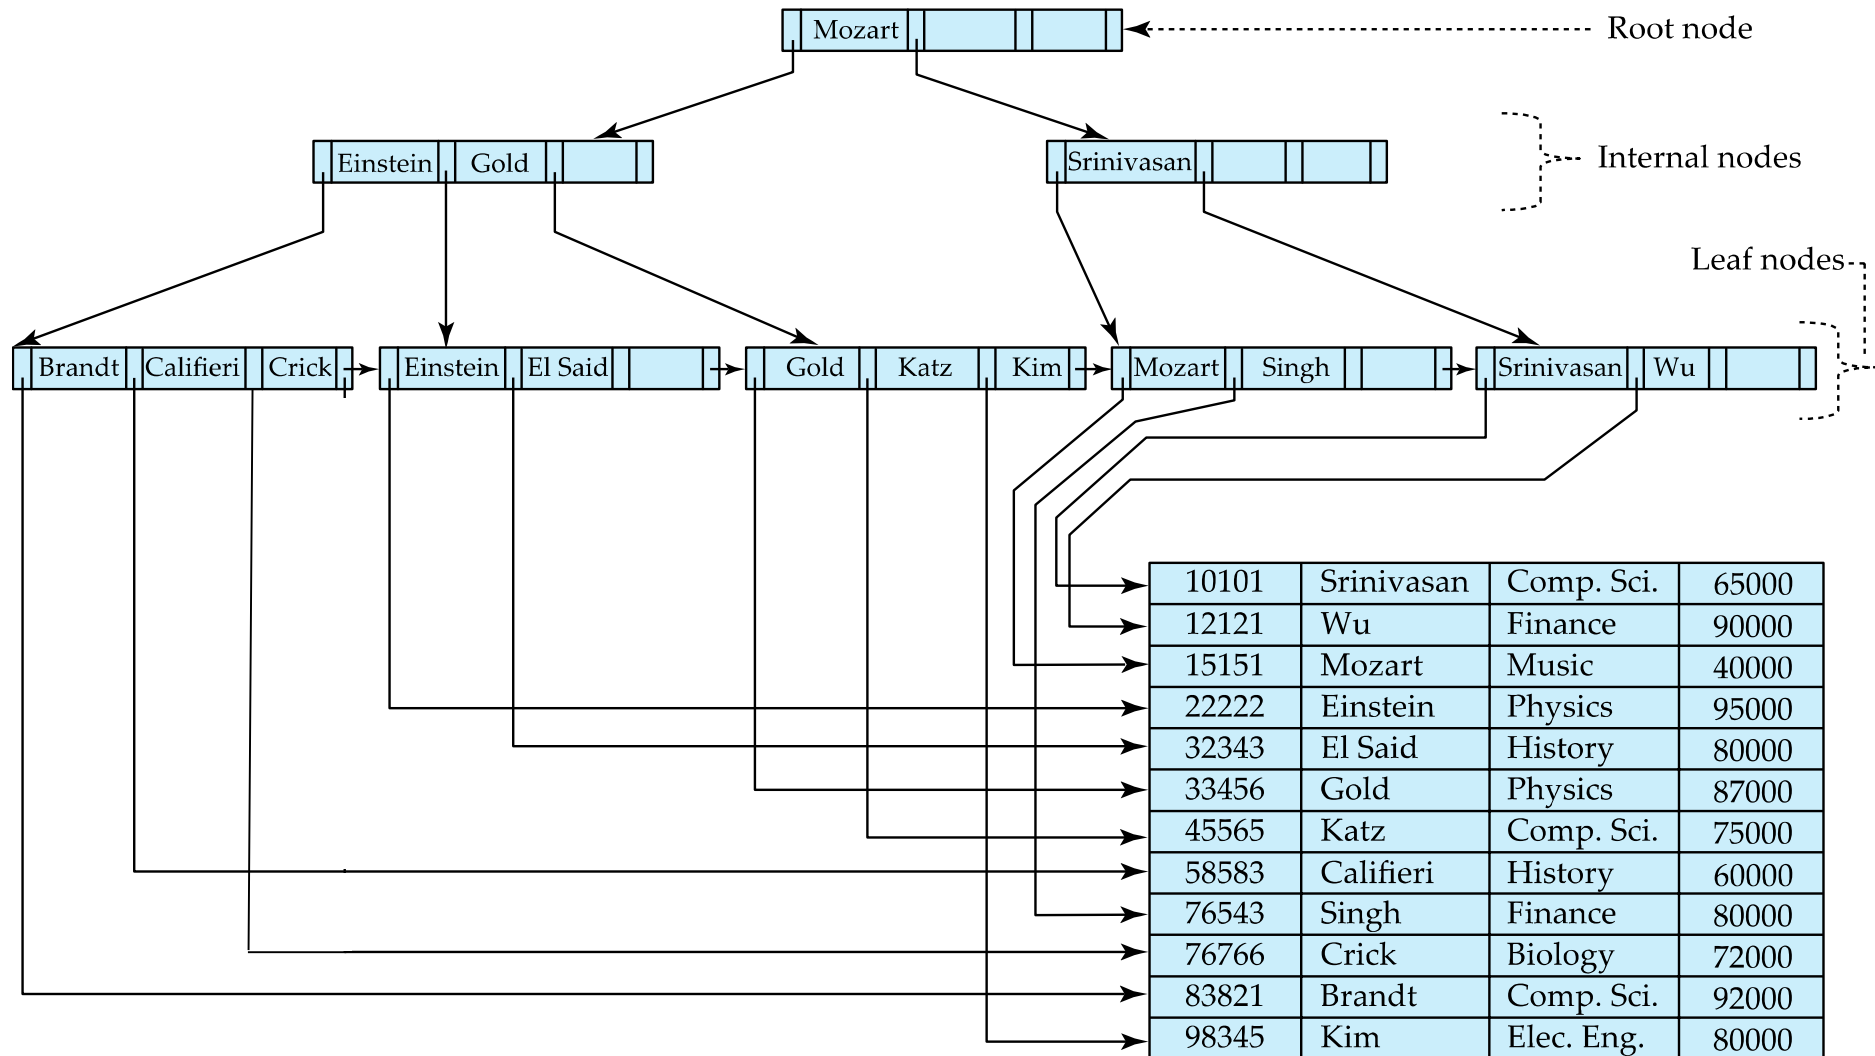


# B<sup>+</sup>-Tree Index Files

- In terms of the number of values or slots, a node that is not a root or a leaf can be slightly less than half-full for  $n$  even, but would be at least half-full for  $n$  odd:
  - For  $n$  even, the maximum number of values is  $(n - 1)$ , and the minimum number of values is  $\lceil n/2 \rceil - 1 = n/2 - 1$ , which gives a minimum fullness of  $(n/2 - 1)/(n - 1) = \frac{1}{2} [(n - 2)/(n - 1)] < \frac{1}{2}$ , which will  $\rightarrow \frac{1}{2}$  as  $n \rightarrow \infty$ .
  - For  $n$  odd, the maximum number of values is  $(n - 1)$ , and the minimum number of values is  $\lceil n/2 \rceil - 1 = (n - 1)/2$  as shown earlier, which gives a minimum fullness of  $[(n - 1)/2]/(n - 1) = \frac{1}{2}$
- **To ensure that the number of slots of a leaf node is at least half-full, instead of using the lower limit  $\lceil n/2 \rceil - 1$ , which for even  $n$  can be less than half-full, we modify it to  $\lceil n/2 \rceil$**
- If the root is not a leaf, it has at least 2 children



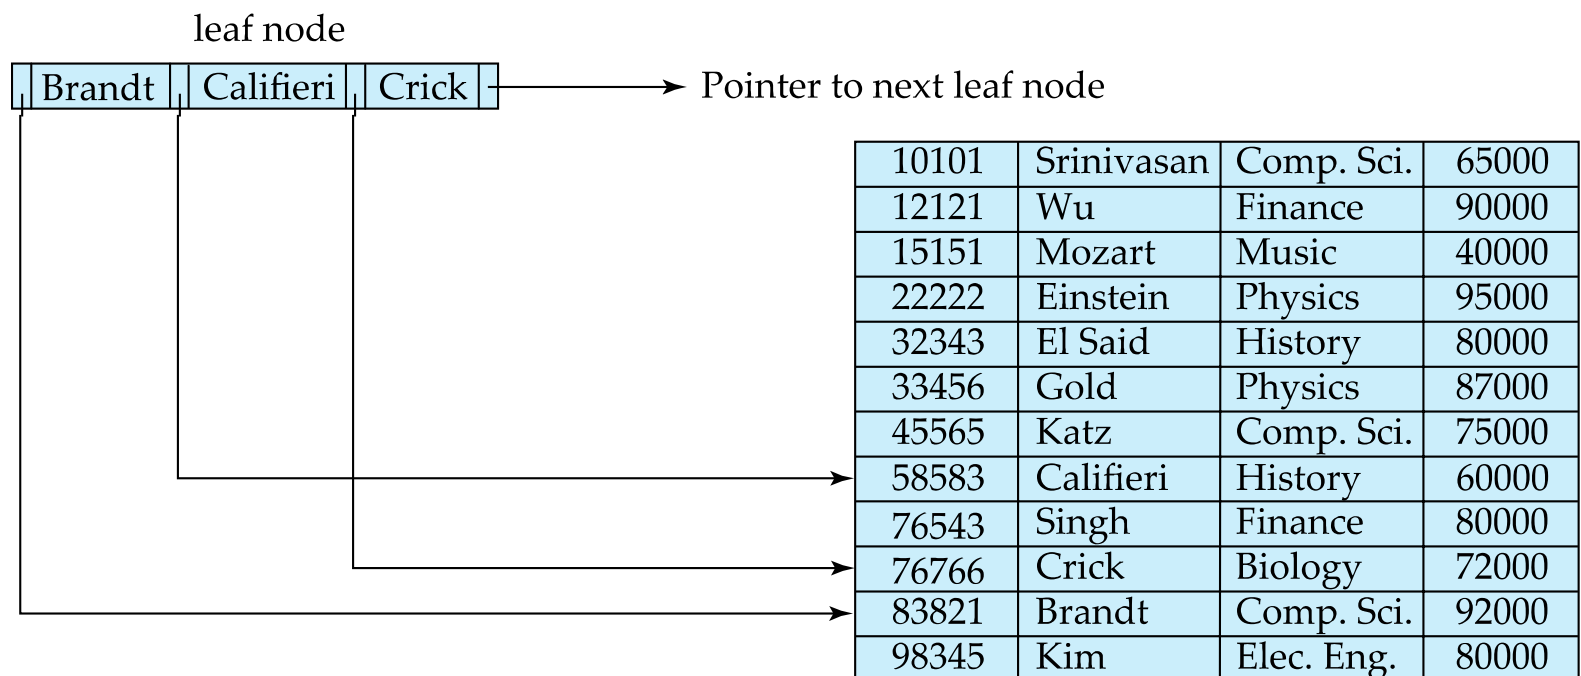
# Example of B<sup>+</sup>-Tree ( $n = 4$ )





# Leaf Nodes in B<sup>+</sup>-Trees ( $n = 4$ )

- For  $i = 1, 2, \dots, n-1$ , pointer  $P_i$  points to a record with search-key value  $K_i$ ,
- If  $L_i, L_j$  are leaf nodes and  $i < j$ ,  $L_i$ 's search-key values are less than or equal to  $L_j$ 's search-key values
- $P_n$  points to next leaf node in search-key order (facilitates sequential processing)








# Example of B<sup>+</sup>-tree ( $n = 6$ )

- B<sup>+</sup>-tree for *instructor* file ( $n = 6$ ), i.e., 6 pointers and 5 slots



- Leaf nodes must have between 3 and 5 values ( $\lceil n/2 \rceil$  and  $n-1$ , with  $n = 6$ ).
- Non-leaf nodes other than root must have between 3 and 6 children ( $\lceil n/2 \rceil$  and  $n$  children)  $\Rightarrow$  between 2 and 5 values
  - In terms of the values, this is less than half-full (i.e.,  $2/5 = 40\%$  full)
- If, however,  $n = 7$ , then the non-leaf nodes will have between 4 and 7 children ( $\lceil n/2 \rceil$  and  $n$  children)  $\Rightarrow$  between 3 and 6 values which is at least half-full
- Root must have at least 2 children



# Performance of B<sup>+</sup>-trees

- For a B<sup>+</sup>-tree, where each node contains between  $m = \lceil n/2 \rceil$  and  $n$  children (assuming the root behaves like any other node):
  - Number of nodes at Level 1 =  $m^0$  (root)
  - Minimum number of nodes at Level 2 =  $m^1$
  - Minimum number of nodes at Level 3 =  $m^2$
  - Minimum number of nodes at Level 4 =  $m^3$
  - Minimum number of nodes at Level  $h = m^{h-1}$
- Each leaf node will hold roughly at least  $m$  values, so that assuming the height of the tree is  $h$ , the minimum number of values held by the leaf nodes is approximately  $m \times m^{h-1} = m^h$
- If there are  $K$  search-key values in the file, we have
$$m^h = K$$
giving  $h = \lceil \log_{\lceil n/2 \rceil}(K) \rceil$
- By assuming minimum storage utilization of each node, we have the maximum number of nodes for the tree and hence maximum height for the tree
- Thus, if there are  $K$  search-key values in the file, the tree height should be no more than  $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$  (i.e. we are assuming all nodes are minimally full; if they are not minimally full, then the height should be less)



# Performance of B<sup>+</sup>-Trees

- If there are  $K$  search-key values in the file, the height of the tree is no more than  $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$ .
- A node is generally the same size as a disk block, typically 4 kilobytes
  - Assuming 40 bytes per index entry,  $n$  is typically around 100
- With 1 million search key values and  $n = 100$ 
  - at most  $\log_{50}(1,000,000) = 3.53 \approx 4$  nodes are accessed in a lookup traversal from root to leaf
- Contrast this with a balanced binary tree with 1 million search key values — around 20 nodes ( $\log_2(1,000,000) = 19.93 \approx 20$ ) are accessed in a lookup
  - above difference is significant since every node access may need a disk I/O, costing around 20 milliseconds



# Average Storage Utilization of B-Tree

- Let  $K$  be the total number of items in the tree,  
 $n$  be the maximum capacity of a node (i.e., a node can hold between  $n/2$  and  $n$  items)  
 $N$  be the random number of nodes in the tree  
 $\rho$  be the random storage utilization of the tree  
 $f$  be the minimum fullness factor
  - $f = 1/2$  for the standard B-Tree,
  - $f = 2/3$  for the B\*-Tree
- The total storage capacity of the (random) tree is  $Nn$
- The storage utilization is the total number of items divided by the total storage capacity of all the nodes, i.e.,
$$\rho = K/(Nn)$$
- Now a minimum number of nodes would result if all nodes are full, which is  $K/n$
- Likewise, a maximum number of nodes would result if all nodes are half-full, which is  $K/(1/2n) = 2K/n$



# Average Storage Utilization of B-Tree

- For random insertion in which every configuration in the above node range is equally likely, the distribution of  $\mathbf{N}$  may be approximated by the continuous uniform distribution over the interval  $[K/n, 2K/n]$  of length  $K/n$ , with height  $n/K$ .

- That is we have

$$\mathbf{N} \sim \mathcal{U}(K/n, 2K/n)$$

where  $\mathcal{U}$  signifies the uniform distribution

- Thus, we have, approximately,

$$E(\rho) = E(K[\mathbf{N}n]) = (K/n) E(1/\mathbf{N})$$

$$= \left(\frac{K}{n}\right) \times \left(\frac{n}{K}\right) \int_{\frac{K}{n}}^{\frac{2K}{n}} \left(\frac{1}{t}\right) dt = \ln\left(\frac{2K}{n}\right) - \ln\left(\frac{K}{n}\right) = \ln 2 = 69.3\%$$



# Average Storage Utilization for the General Case

- For the general case with arbitrary minimum fullness  $f$ , the distribution of  $\mathbf{N}$  may be approximated by the uniform distribution over the interval  $[K/n, K/(nf)]$  of length  $Kf'/(nf)$ , with height  $nf/(Kf')$ , where  $f' = 1 - f$ .
- That is we have

$$\mathbf{N} \sim \mathcal{U}(K/n, K/nf)$$

- Thus, we get, approximately,

$$\begin{aligned} E(\rho) &= E(K[\mathbf{N}n]) = (K/n) E(1/\mathbf{N}) \\ &= \left(\frac{K}{n}\right) \times \left(\frac{nf}{Kf'}\right) \int_{\frac{K}{n}}^{\frac{K}{nf}} \left(\frac{1}{t}\right) dt = \frac{f}{f'} \ln \frac{1}{f} \end{aligned}$$

- For the B\*-Tree,  $f = 2/3$ , and substituting this value into the above, we get

$$E(\rho) = 2 \times \ln(3/2) \approx 81\%$$



# Variance of the Storage Utilization

The variance can be shown to be:

$$\sigma_f^2 = f - \left( \frac{f}{f'} \right)^2 \left[ \ln \left( \frac{1}{f} \right) \right]^2$$

The standard deviation of storage utilization of the B-Tree is approximately 14%, and that of the B\*-Tree is 9.4%.



# Updates on B<sup>+</sup>-Trees: Insertion

Let

1.  $Pr$  be pointer to the record, and let
2.  $V$  be the search key value of the record

Find the leaf node in which the search-key value would appear


1. If there is room in the leaf node, insert  $(V, Pr)$  pair in the leaf node
2. Otherwise, split the node





# Updates on B<sup>+</sup>-Trees: Insertion

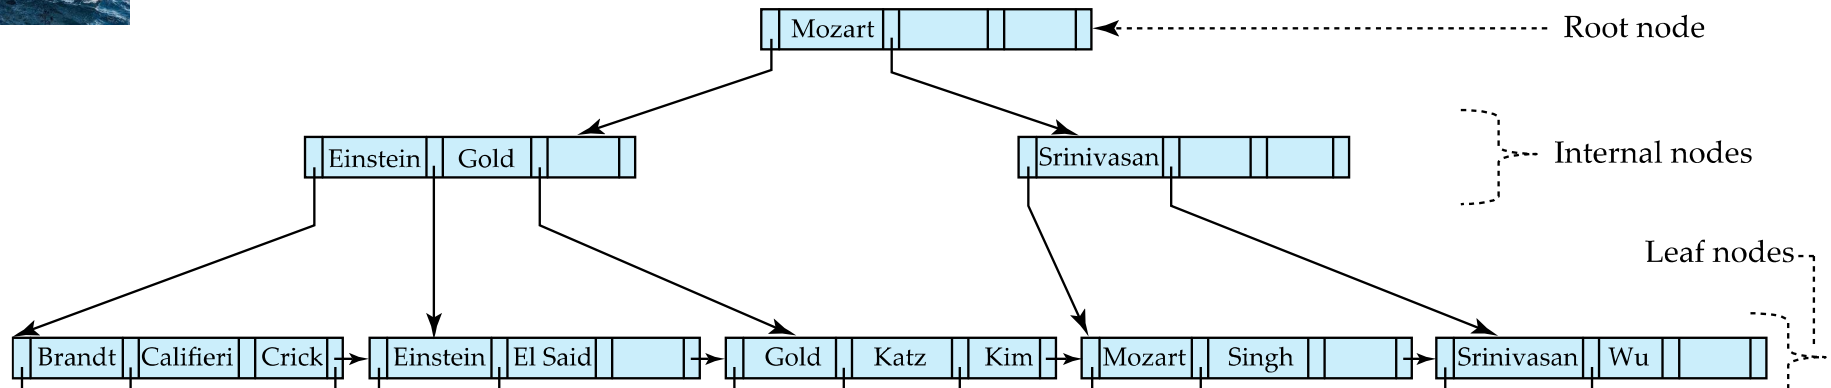
- Splitting a leaf node:
  - take the  $n$  (search-key value, pointer) pairs (including the one being inserted) in sorted order. Place the first  $\lceil n/2 \rceil$  in the original node, and the rest in a new node
  - let the new node be  $p$ , and let  $k$  be the least key value in  $p$ . Insert  $(k, p)$  in the parent of the node being split
  - If the parent is full, split it and **propagate** the split further up
- Splitting of nodes proceeds upwards till a node that is not full is found
  - In the worst case the root node may be split increasing the height of the tree by 1



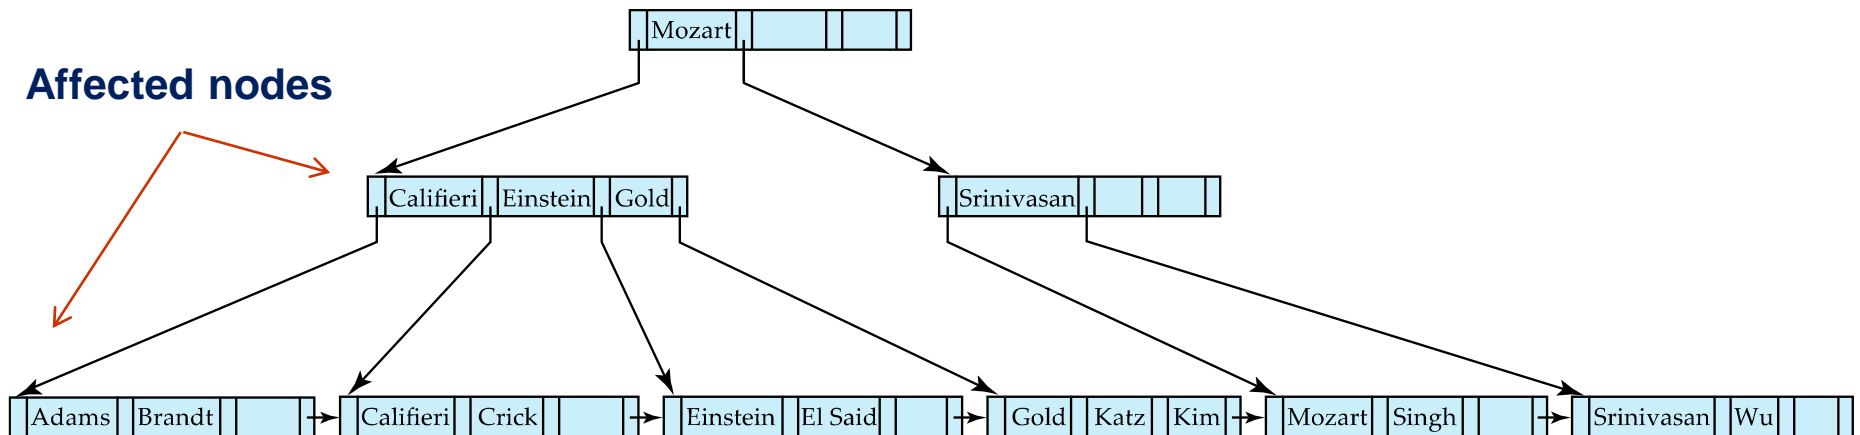
Result of splitting node containing Brandt, Califieri and Crick on inserting Adams  
Next step: insert entry with (Califieri, pointer-to-new-node) into parent



# B<sup>+</sup>-Tree Insertion ( $n = 4$ )

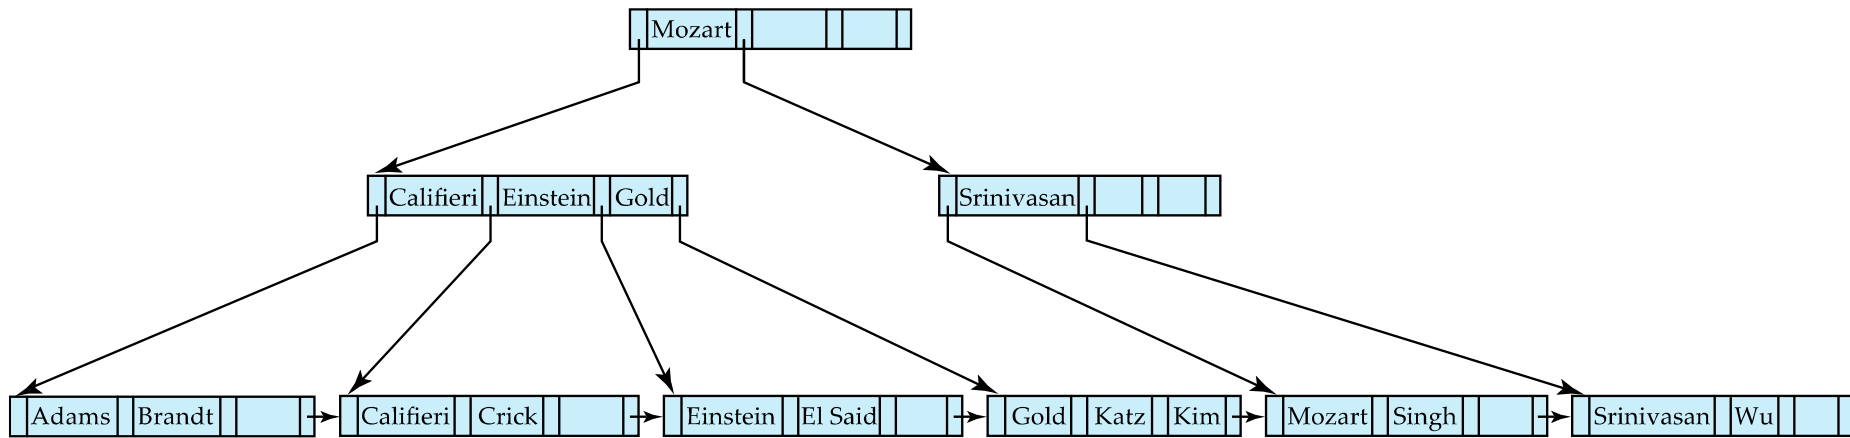


## B<sup>+</sup>-Tree before and after insertion of "Adams"

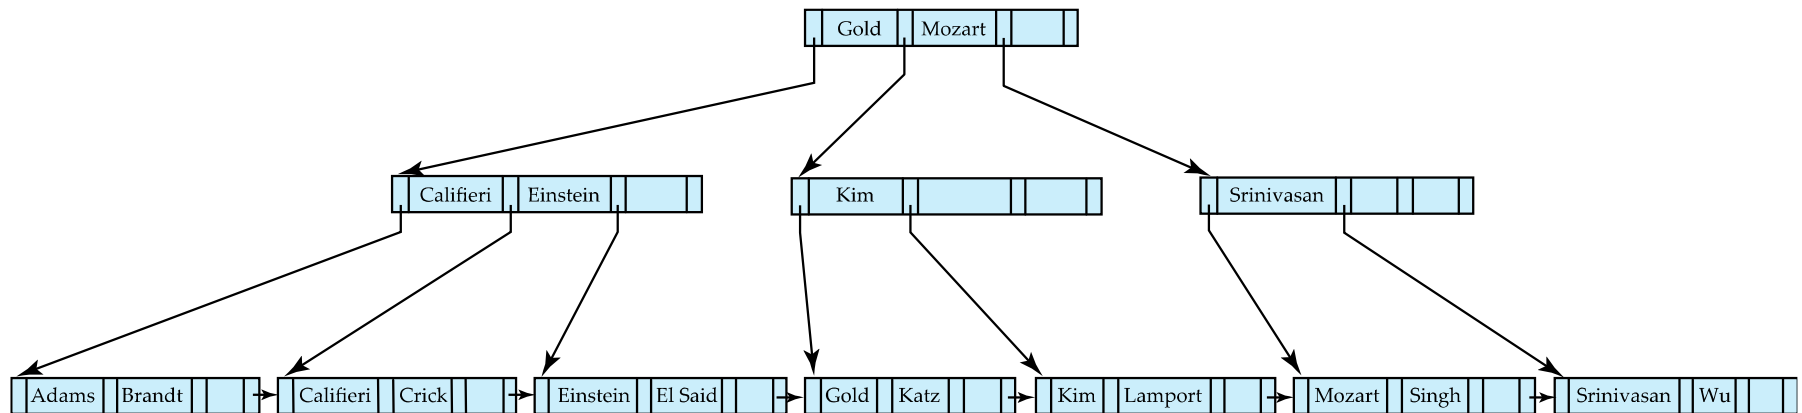




# B<sup>+</sup>-Tree Insertion ( $n = 4$ )

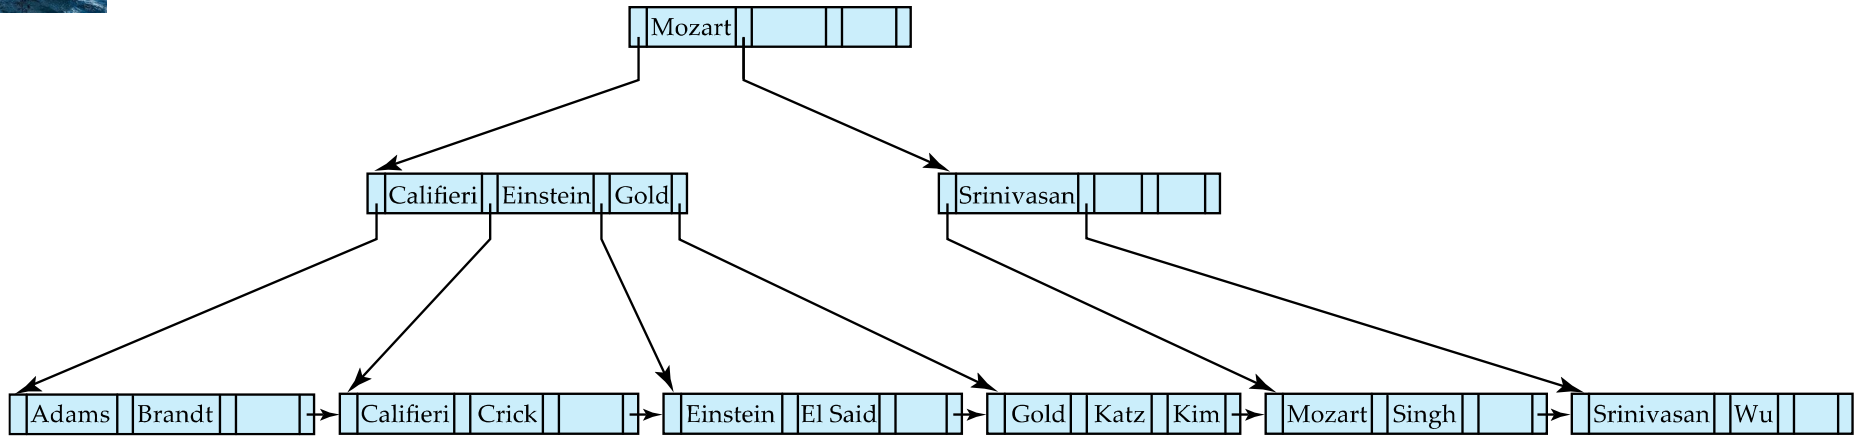


## B<sup>+</sup>-Tree before and after insertion of “Lamport”

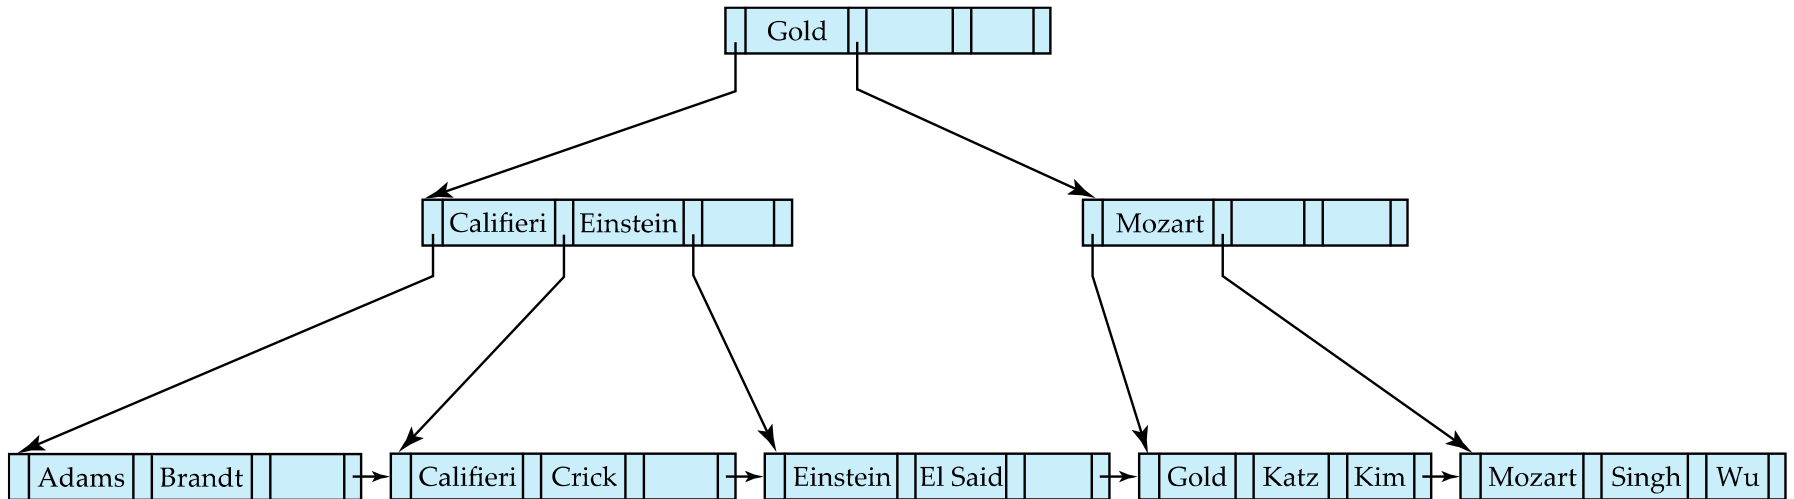




# Examples of B<sup>+</sup>-Tree Deletion ( $n = 4$ )



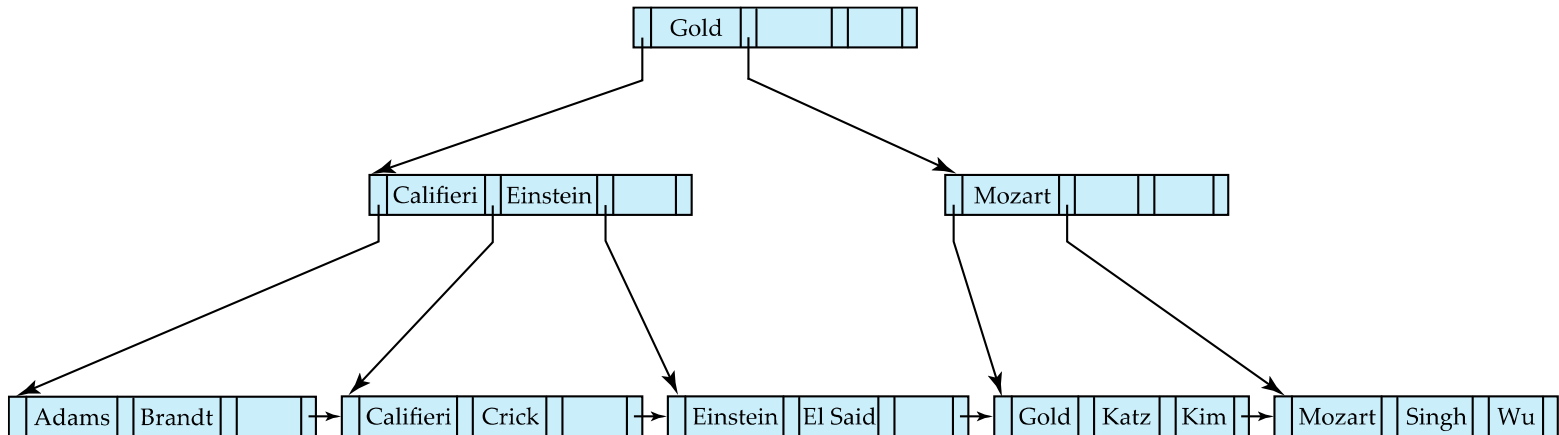
Before and after deleting “Srinivasan”




- Deleting “Srinivasan” causes **merging** of under-full leaves



# Examples of B<sup>+</sup>-Tree Deletion ( $n = 4$ )



Before and after deleting “Singh” and “Wu”



- Leaf containing Singh and Wu became under-full, and **borrowed a value** Kim from its left sibling
- Search-key value in the parent changes as a result



# Updates on B<sup>+</sup>-Trees: Deletion

Assume record already deleted from file. Let  $V$  be the search key value of the record, and  $Pr$  be the pointer to the record

- Remove  $(Pr, V)$  from the leaf node
- If the node has too few entries due to the removal, and the entries in the node and a sibling fit into a single node, then **merge siblings**:
  - Insert all the search-key values in the two nodes into a single node (the one on the left), and delete the other node
- If the node has too few entries due to the removal, but the entries in the node and a sibling do not fit into a single node, then **redistribute pointers**:
  - Redistribute the pointers between the node and a sibling such that both have more than the minimum number of entries
- Update the parent node
- The node deletions may cascade upwards



# Variations

- B<sup>+</sup>-Tree File Organization
  - Stores actual records in leaf node, not just pointers
- B-Tree Index Files
  - Points to actual records in non-leaf nodes
- B<sup>\*</sup>-Tree
  - Can vary minimum fullness factor  $f$  to be different from  $f = \frac{1}{2}$ ,
  - for  $f = \frac{2}{3}$ , the corresponding tree is called a B<sup>\*</sup>-Tree